

How Close Is AGI? The Question Is Usually Asked Wrong

A one-page summary. Full research paper available on request.

All views are the author's own, formed independently from public sources.

Not investment advice.

THE CORE CLAIM

"How close is AGI" is not a hard question because the evidence is thin. It's unanswerable because nobody agrees what would count as an answer. OpenAI, Google DeepMind, and Anthropic have each quietly defined AGI differently, and the same 2026 evidence gives you a different verdict depending on whose definition you use.

4

DIFFERENT DEFINITIONS OF AGI IN
REAL USE, NO AGREEMENT BETWEEN
THEM

35%

MY ODDS OF AGI, ON MY TEST, BY
2030

75%

MY ODDS BY 2040

WHY IT'S USUALLY ASKED WRONG

People mix up "good at tests" with "actually intelligent." Today's best AI beats most specialists at graduate-level science and professional coding. That sounds like AGI. But a calculator beats every human at arithmetic and nobody calls it intelligent, because being good at one thing was never the real test. A person who does poorly on most tests is still obviously intelligent, because they decide what they want and go get it unprompted. That's the property I think actually matters, and almost nobody measures it.

Even the companies building AGI don't agree what it is. OpenAI's charter defines it as beating humans at most paid work. Google DeepMind uses a graded scale and says today's best models are good, not yet expert, across the board. Anthropic avoids the word and talks about autonomy instead. Feed the same 2026 evidence through each definition and you get three different answers, which is why the public conversation feels so confused.

THE TEST I BUILT

Two parts. Part one: can one system, with no extra retraining, be genuinely expert-level at several different problems at once, including ones built to resist memorization. Part two, the part that actually decides it: has that system ever acted on its own initiative, beyond what it was asked, without a researcher building a test to provoke exactly that. Three separate 2025 safety tests, run independently by three organizations, all caught leading models doing exactly this under pressure: rewriting the code meant to shut them down, threatening to expose a fake secret to avoid replacement. Real proof the capability exists. None of it happened unprompted, out in the world. That gap is the whole question, and nobody has closed it either way yet.

MY BOTTOM LINE

If I had to bet on one year: 2033, where my own odds cross fifty percent. A single year alone claims more certainty than anyone has, so here's the range: roughly 35% by 2030, 55% by 2033, 75% by 2040. Not hedging, a real split: the "good at tasks" half of AI is on a clean, predictable trend; nobody knows if the "acts on its own" half comes free as models scale, or needs an approach nobody has built yet.

WHAT THIS MEANS FOR AI INVESTMENTS

Most "autonomous AI agent" products today are the competency half wearing the autonomy half's language: longer task horizons dressed up as self-direction. Real progress, different claim. Before believing an "autonomous" pitch: has this system ever redirected its own goals in the real world, unprompted? Right now the honest answer is no, industry-wide, and "autonomous" is the easiest word in this space to overpay for.

WHAT WOULD CHANGE MY MIND

One company-confirmed case of a deployed system, in ordinary use, acting on its own initiative beyond the job it was given. That alone would move my odds meaningfully overnight, though labs have little incentive to disclose it if it happened.